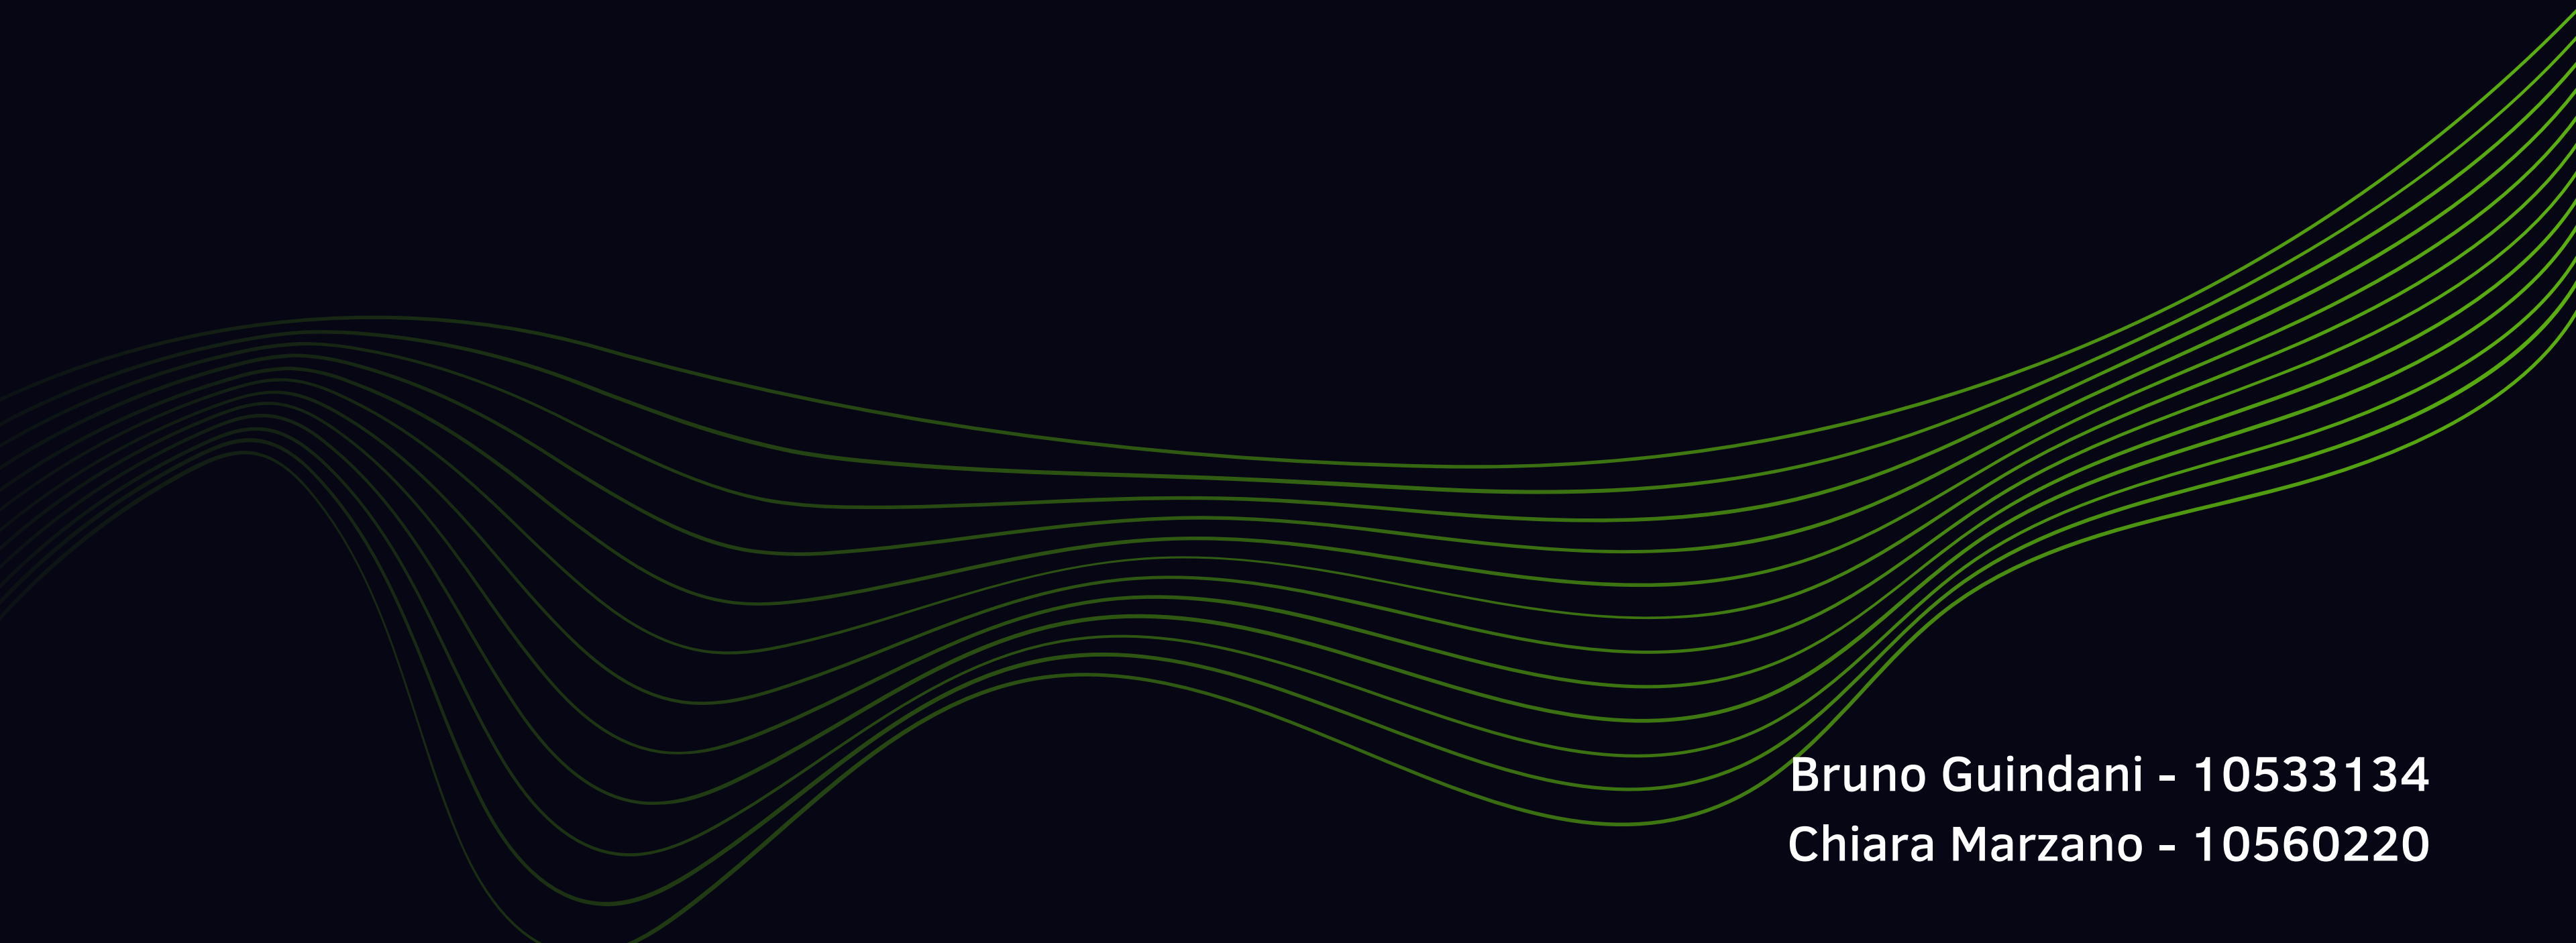


Parallel Computing Presentation

01

A series of approximately 15 thin, light green wavy lines that flow from the left side of the slide towards the right, creating a sense of motion and depth.

Bruno Guindani - 10533134
Chiara Marzano - 10560220

Agenda

02

•

ORIGINAL PROBLEM
Assessment of the original
problem, profiling and
measurements

•

**TECHNOLOGY
CHOICES**
CUDA for GPGPU

•

IMPLEMENTATION
Where and how we
implemented
parallelism

•

PERFORMANCE
Evaluation of our
solution

•

**NUMERICAL
RESULTS**
Comparison with
the original

The background is a dark blue gradient with numerous thin, wavy, light blue lines that create a sense of movement and depth. A single small white dot is located on the left side of the image.

ORIGINAL PROBLEM

03



Original Problem

Interaction of particles in space

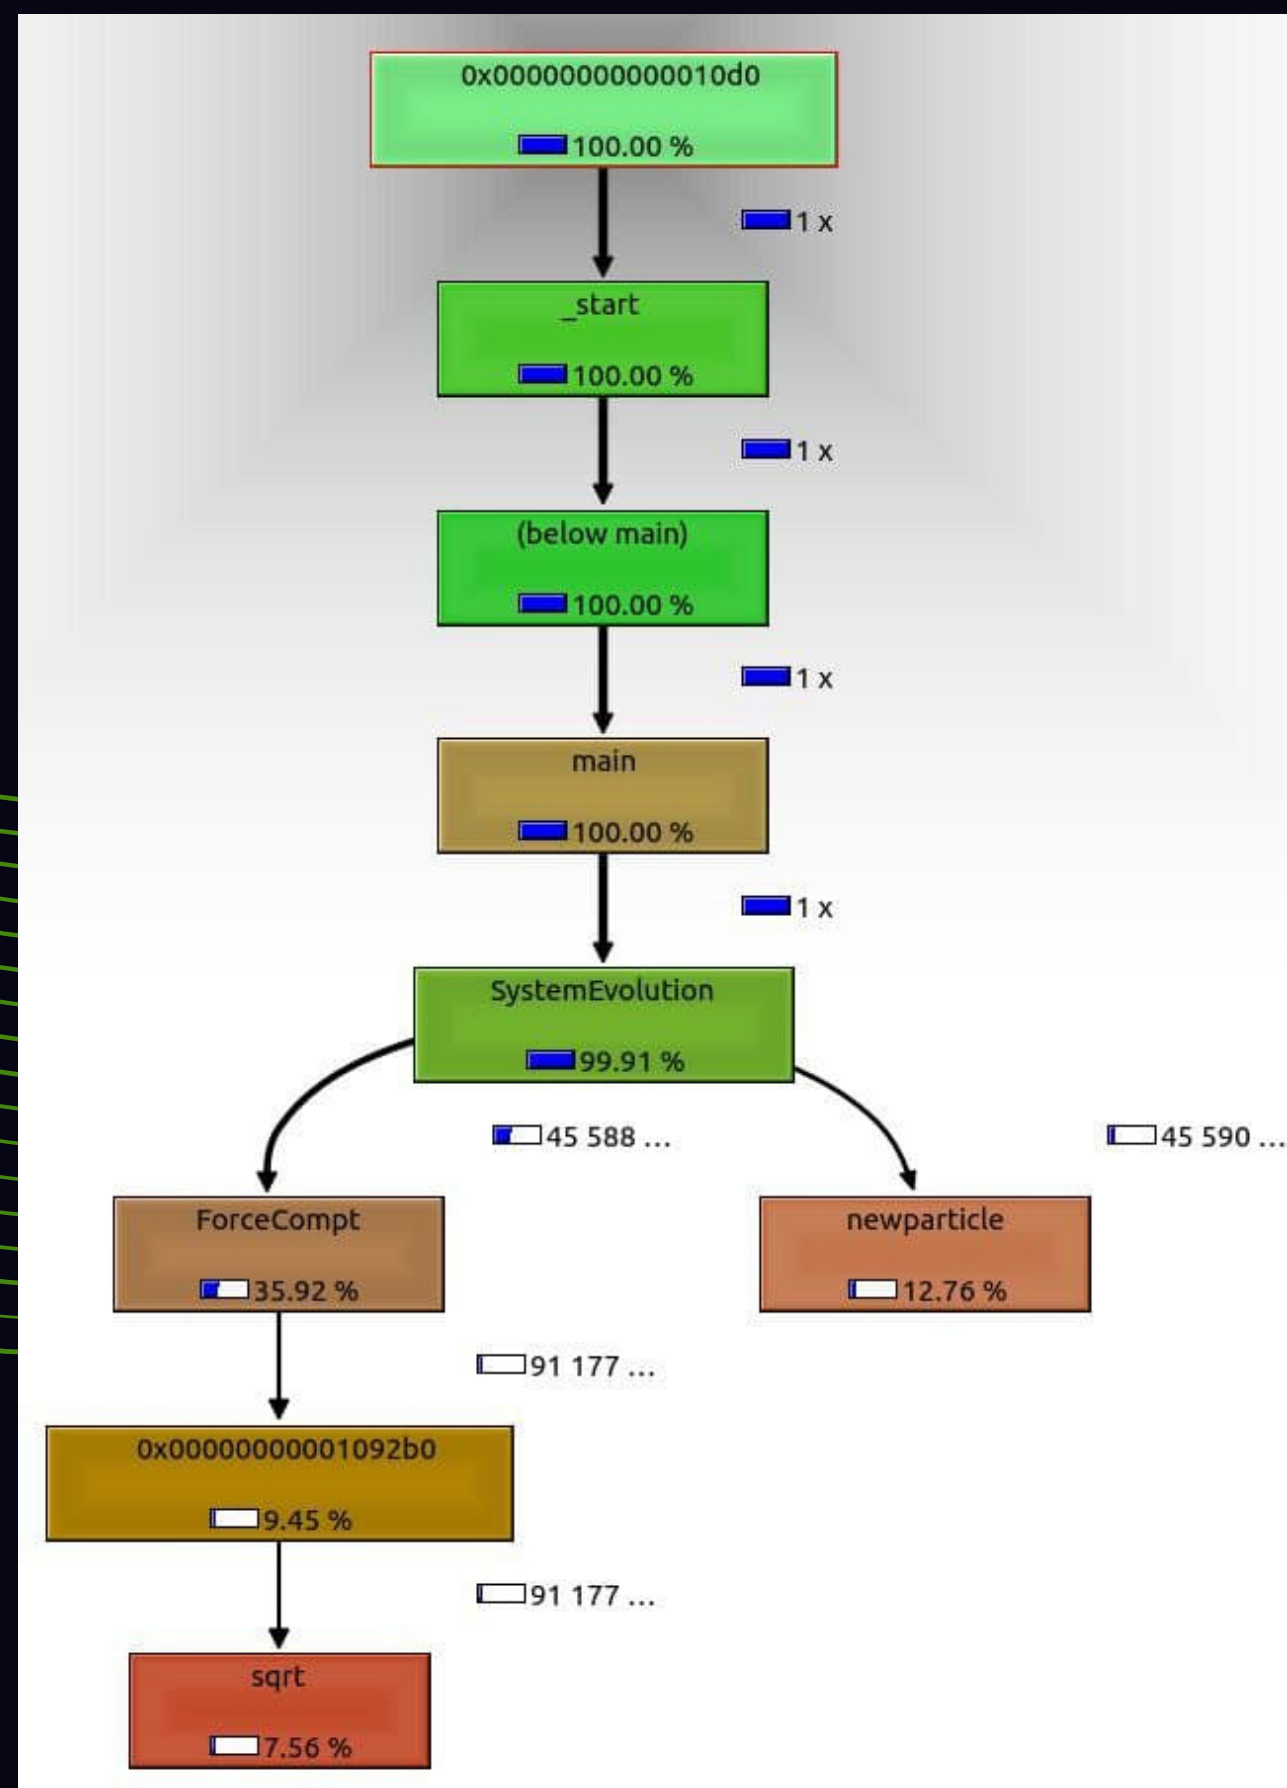
GENERATING FIELD

PARTICLE GENERATION

SYSTEM EVOLUTION




04



Performance Assessment

using KCacheGrind

OUR SOLUTION



Technology choice

CUDA for GPGPU

FINE-GRAINED CONTROL

EDUCATIONAL CHALLENGE

HIGHLY SUITABLE
APPLICATION DOMAIN

07

Implementation

Selecting functions to parallelize

Our first focus was deciding what to execute GPU-side and CPU-side. The decision was mainly based on:

- bottlenecks
- highly-parallelizable and concurrency-free operations
- educational impact

08

Implementation


Techniques

- mono and bidimensional stride
- shared memory

09

Implementation: stride technique

```
idx = blockIdx.x * blockDim.x + threadIdx.x;  
stridex = gridDim.x * blockDim.x;
```





0	1	2
3	4	5
6	7	8

9	10	11
12	13	14
15	16	17

18	19	20
21	22	23
24	25	26

27	28	29
30	31	32
33	34	35


+36

second stride

36	37	38
39	40	41
42	43	44

45	46	47
48	49	50
51	52	53

54	55	56

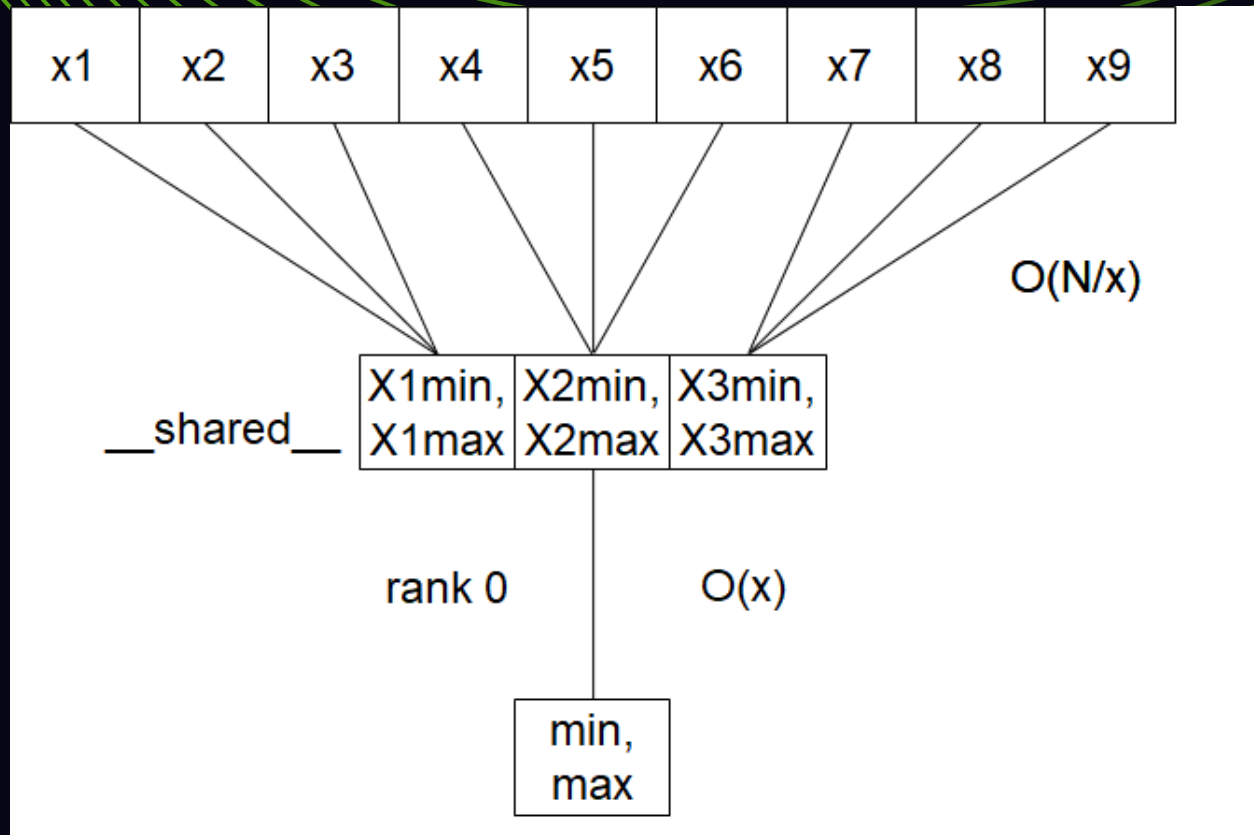


11

Generating Field

- Computes the generating field grid values in all points of the defined space
- Deterministic generation
- $EX * EY$ independent iterations
- Bidimensional stride technique
- Bidimensional threads per block: 32, 32
- Bidimensional number of blocks: $32 * \text{num_SMs}$, $32 * \text{num_SMs}$

Minimum and maximum



- Joint MinMaxInt() and MinMaxDouble() functions
- Single block, shared memory
- Optimal number of threads: $\operatorname{argmin} f(x) = x + N/x$
→ $x = \sqrt{N}$ (capped at 1024)

13

System Instant Evolution

- Executes a single evolution iteration
- NP independent iterations
- Monodimensional stride technique
- Monodimensional threads per block: 1024
- Monodimensional number of blocks: $32 * \text{num_SMs}$

14

InitializeEmpty Grid

- Initializes empty particle grid
- EX * EY independent iterations
- Bidimensional stride technique

15

Compt Population

- Computes effects of forces on particles
- NP independent iterations
- Monodimensional stride technique

16

Performance

data collected over 10 runs, in microseconds

1200 * 1000

Sequential

- average: 234'457'848
- standard deviation: 2'478'229

Parallel

- average: 11'110'354
- standard deviation: 110'340

Speedup: 20x

2400 * 2000

Sequential

- average: 895'111'148
- standard deviation: 26'741'978

Parallel

- average: 42'099'243
- standard deviation: 195'185

Speedup: 20x

4800 * 4000

Sequential (one iteration)

- average: 373'628'321
- standard deviation: 3'001'303

Parallel (one iteration)

- average 2'189'173
- standard deviation: 20068

Speedup: 170x


17



18

Numerical Results

- Sometimes, the populations generated by the sequential and parallel programs are identical except for a single particle



particle number 214
sequential weight: 1650
parallel weight: 1950

19

THANK YOU FOR
YOUR
ATTENTION

Bruno Guindani - 10533134
Chiara Marzano - 10560220

20

Compiling and running

REPOSITORY

Can be cloned at <https://github.com/ChiaraMarzano/ParComProject>

COMPILING

```
nvcc -arch=sm_70 -o particles.x particles_c.cu
```

RUNNING

```
./particles.x
```